

Research Proposal

Design-Based Sensitivity Analysis for Matched Observational Studies under Hidden Bias, Population Shift, and Model-Generated Measurements

1. Summary

Matched observational studies are the standard design for estimating treatment effects when randomization is impossible, and their credibility rests on one object: the distribution of treatment assignments within matched sets. Classical practice takes that distribution to be uniform and asks, through a sensitivity analysis, how far it could depart from uniformity before a conclusion is overturned [1]. My current work with Prof. Xinran Li [2] shows that the anchor should instead be the fitted propensity odds, that departures should be bounded on the odds scale by a parameter Λ , and that a lower-envelope statistic gives studentized tests of Neyman’s weak null that are valid over the whole band of admissible laws. This proposal generalizes that principle, an **anchor**, a **bounded departure**, and a **guarantee over the band**, to the three ways in which the data of a modern matched study depart from the ideal at once: the assignment law may depend on unmeasured covariates (**hidden bias**); the matched sample is rarely the population in which the conclusion will be used (**population shift**); and covariates and outcomes are increasingly produced by language models from text rather than measured (**model-generated measurements**). Each departure is treated as a perturbation of a law, anchored at an estimate, bounded by an interpretable parameter, and carried through randomization inference rather than assumed away. The three aims are (1) design sensitivity and optimal contrasts for covariate-adaptive full matching; (2) joint sensitivity analysis for hidden bias and transport to a shifted target population; (3) design-based inference with model-generated covariates and outcomes and a small human-verified calibration sample. Each aim delivers theory, open-source software, and an application in clinical or science-of-science data.

2. Background and preliminary results

2.1 Setting. There are K matched sets; set k contains n_k units with treatment indicators Z_{kj} , potential outcomes $Y_{kj}(1), Y_{kj}(0)$, and observed outcome $Y_{kj} = Z_{kj}Y_{kj}(1) + (1 - Z_{kj})Y_{kj}(0)$. The target is Neyman’s one-sided weak null $H_0(\tau_0) : \bar{\tau}_K \leq \tau_0$ on the matched-sample average effect, which restricts only the average and imputes no missing potential outcome. Randomization inference for weak nulls requires studentization to remain valid under heterogeneous effects [3, 4]; Rosenbaum’s model allows the treatment odds of two units in a set to differ by at most Γ and reports the largest Γ at which a conclusion survives [1]. Inexact matching leaves residual propensity discrepancies that invalidate the uniform reference law [5], and covariate-adaptive randomization inference replaces it with probabilities derived from fitted propensity odds [6].

2.2 My results. In [2] we combine the two. With baseline odds r_{k1}, r_{k2} in pair k , the anchor is $\theta_k = r_{k1}/(r_{k1} + r_{k2})$; a unit-level hidden-bias bound Γ_u induces, after conditioning on one treated unit per pair, the pair-level parameter $\Lambda = \Gamma_u^2$ and the interval $q_k \in [L_k(\Lambda), U_k(\Lambda)]$ with $L_k = \theta_k/\{\theta_k + (1 - \theta_k)\Lambda\}$ and $U_k = \theta_k/\{\theta_k + (1 - \theta_k)/\Lambda\}$. The proposed statistic is an asymmetric Horvitz–Thompson lower envelope: a positive observed contrast is divided by the largest feasible probability of observing it and a negative one by the smallest. We prove that the envelope score is mean-controlled and pointwise sharp within a zero-normalized observed-cell class, that it reduces to the known-propensity Horvitz–Thompson statistic at $\Lambda = 1$, and that its studentized average gives asymptotically valid tests, with a randomization-calibrated critical value, an honest inflated-endpoint version for estimated propensities, and a validity theorem for the plug-in test under fast propensity estimation. The construction extends to full matching [7] in closed form: a level-free difference

contrast depends on the assignment vector only through the realized special unit's coordinate, whose marginal range is again a scalar interval, so the entire theory carries over. Within full matching, the simple within-set mean is one member of a family of balanced, pretreatment matrix-weighted contrasts W that share every guarantee, and a propensity-anchored local member recovers substantial power in large heterogeneous sets. In the right heart catheterization study the covariate-adaptive analysis reports a sensitivity value of $\Lambda = 1.40$ against 1.30 for the uniform analysis. Two features of this work carry into the proposal: the anchor is a function of estimated propensities, so any error in the covariates displaces it (Aim 3), and the choice of W and of the matched structure is a free design decision whose consequences for power are not yet understood (Aim 1).

2.3 Related work and gap. Design sensitivity [8] and split-sample design choice [9] exist for uniform reference laws and sharp nulls. Transport of causal conclusions to a target population uses estimated density-ratio weights [10, 11], with sensitivity analyses for unobserved effect modifiers [12, 13] and for misspecified weights [14], and average-case sensitivity analysis offers a less conservative alternative to worst-case bounds [20]; none of these is design-based in the randomization sense or combined with hidden bias in assignment. Prediction-powered inference [15–17] and design-based supervised learning [18] correct model-generated labels with a human-labelled subsample, but for population means and regressions, not for matched designs, sensitivity analysis, or the covariate side. That is the gap this proposal fills.

3. Aim 1: Design sensitivity and optimal contrasts for covariate-adaptive full matching

Problem. Our theory shows that in full matching a whole family of weighted contrasts is valid, but says nothing about which member to use or which matched structure (pairs, fixed ratio, full matching with bounded set sizes) to build. Rosenbaum's design sensitivity $\tilde{\Gamma}$, the limiting sensitivity parameter up to which a study can reject as $K \rightarrow \infty$ under a favourable model with a real effect and no hidden bias, is the natural criterion [8], but it has been derived only for uniform reference laws.

Approach. Under a favourable model $Y(0) = h(\text{logit } e(X)) + \varepsilon$, $Y(1) = Y(0) + \tau$, I will derive the design sensitivity of the covariate-adaptive studentized weak-null test as a function of the set-size distribution and of W , and characterize the maximizing W ; the conjecture, supported by our simulations, is that the propensity-anchored local weight is optimal whenever propensity locality is prognostically informative. I will then quantify how anchoring itself changes design sensitivity relative to the uniform law when residual discrepancy is prognostic, which turns the empirical finding in the RHC study into a theorem. Because h and the error scale must be estimated, and estimation requires outcomes, I will follow the split-sample protocol of [9], choosing W and the structure on a planning sample and testing on the remainder, and study the optimal planning fraction when the design is chosen by search over many candidates. Recent work on model-assisted randomization tests that avoid sample splitting by fitting once [19] suggests a route to recovering the lost sample, which I will pursue with Prof. Yao Zhang. The computational deliverable is a solver for the envelope optimization over general set structures.

4. Aim 2: Joint sensitivity analysis for hidden bias and population shift

Setting. The matched study is drawn from a source population S ; the conclusion is needed in a target population T with covariate distribution $p_T(x) \neq p_S(x)$, for example a treatment decision studied in one hospital system and applied in another. The target estimand is the average effect in T . Standard transport reweights each matched set by an estimated density ratio $\hat{w}(x) = \hat{p}_T(x)/\hat{p}_S(x)$ and assumes that effects are exchangeable across populations given X [10, 11].

Problem. Three things go wrong at once. The weights are estimated and may be misspecified; an unmeasured effect modifier may be distributed differently in T ; and hidden bias in the assignment

law interacts with the reweighting, so that a sensitivity analysis performed on either component alone understates the fragility of the conclusion.

Approach. I propose a two-parameter model (Λ, Δ) . As in [2], Λ bounds the departure of the within-set assignment law from its fitted anchor. Δ bounds, on the odds scale, the ratio of the true transport weight to the fitted one, following the marginal sensitivity model for weighting estimators [14]; this single parameter absorbs both weight misspecification and an unmeasured modifier, and can be interpreted as the strength of an unobserved selection variable. Because assignment within sets and selection into S are independent given X , the worst case of the weighted lower envelope over (Λ, Δ) separates across sets and remains in closed form: each set’s envelope score is multiplied by a set-specific extremal weight. Studentizing across sets gives a test whose validity I will prove under estimated weights with a rate condition, extending the plug-in theorem of [2], and whose output is the (Λ, Δ) frontier at which the conclusion is overturned, a two-dimensional analogue of the sensitivity value. Two refinements come from my work on test-time adaptation, where the central lesson was that deciding when *not* to adapt matters as much as the adaptation. First, target regions with poor overlap inflate Δ without limit; I will develop a design-based analogue of overlap trimming [21] that restricts the target to a subpopulation on which the density ratio is bounded, and report an honest estimand for that subpopulation. Second, I will study selective transport, in which the analysis reports a transported conclusion only for target strata where the frontier is stable and abstains elsewhere, with a formal guarantee on the error of the reported strata. This aim connects to Prof. Doudou Zhou’s work on doubly robust transfer inference for model accuracy under covariate shift [22].

5. Aim 3: Matched studies with model-generated covariates and outcomes

Setting. Each unit has structured covariates X , a text document T (a clinical note, an abstract), and possibly a confounder U recorded only in the text. A language model produces an extraction $\hat{U} = f(T)$ or a predicted outcome $\hat{Y} = g(T)$ for every unit; a calibration subsample $\mathcal{C}$, drawn with known inclusion probabilities π_k , receives human-coded values. Practice has moved faster than theory: in most applications the model output is used as if measured without error.

Outcomes. Let $\hat{D}_k$ be the within-set lower-envelope score computed from $\hat{Y}$ and D_k the same score from Y . The rectified statistic $\hat{\tau}_{PP} = \sum_k \hat{D}_k + \sum_{k \in \mathcal{C}} \pi_k^{-1} (D_k - \hat{D}_k)$ has, conditional on the assignment, the same expectation as the fully labelled statistic *whatever* the model’s error structure, even when the error depends on treatment through the text. Hence its worst case over Λ is attained at the same extremal law and the envelope machinery applies unchanged; validity never rests on the model being accurate. What is new is the variance, which has an assignment component and a labelling component. I will derive a conservative estimator for the sum, prove asymptotic validity of the studentized weak-null test for every Λ , show that with a tuned weight on the prediction term [16] the test is never less powerful than the test on $\mathcal{C}$ alone, and study the allocation of a fixed labelling budget across sets, for instance toward sets where the model’s predictions are least stable across prompts. When no human labels can be obtained I will develop a (Λ, δ) analysis in which δ bounds the within-set difference in mean model error between treated and control units.

Covariates. Matching and propensity fitting on $(X, \hat{U})$ leaves residual imbalance in U and displaces the anchor. If the logit propensity is Lipschitz in U with constant L , the within-set log odds ratio departs from the anchor by at most $L(|U_{kj} - \hat{U}_{kj}| + |U_{kl} - \hat{U}_{kl}|)$. The extraction error is observed on $\mathcal{C}$ and L is estimated there, so a high quantile of this bound is estimable, giving an extraction-induced parameter Λ_{ext} that, unlike the classical Λ , is identified, and that combines multiplicatively with it on the odds scale. The report is then a function of the unidentified Λ with Λ_{ext} estimated and its sampling uncertainty carried through the studentized test; the technical

core is the lower envelope with an anchor that has estimation error of its own, and the target is validity when $|\mathcal{C}|$ grows more slowly than K . Two ideas from Prof. Dongming Huang’s work enter here: conditional inference via exchangeable sampling, which calibrates the test conditional on the calibration sample by resampling extraction errors exchangeably [23, 24], and catalytic priors, which stabilize the propensity fit in small calibration samples with synthetic data without sacrificing validity [25]. Post-treatment leakage, where the text mentions treatment, will be audited by comparing propensity fits from time-stamped pre-treatment text with fits from the full text.

Abstention. My work on scientific world models showed that a predictor should return *unsupported* rather than a number when its evidence is seed-noisy or context-dependent, and that doing so raised accuracy on the reported cases from 71% to 100%. I will formalize the analogue here: when calibration diagnostics show that the model’s error is unstable across prompts or strata, the analysis abstains on the affected sets, with error control for the reported conclusions in the spirit of conformal and selective inference [26]. Prof. Zhou’s work on conditional evaluation of language models with cheap auxiliary signals [27] and on inference with surrogate-derived models under covariate shift [28] is directly relevant and links Aims 2 and 3.

6. Evaluation and applications

Semi-synthetic studies will attach real text and real model outputs to simulated assignments with controlled hidden bias and controlled shift, so that Λ , Δ , Λ_{ext} , and the labelling error are known, and will check size and power across extraction quality, calibration size, and set structure. The right heart catheterization data, which contain no text, serve as the check that every method reduces to [2] when models and shift play no role. Two applications are planned. The first is clinical: MIMIC-IV and its notes, with severity and comorbidity extracted from admission notes as model-generated covariates, an ICU treatment decision in the structured data as the exposure, and transport across care units as the shifted target. The second is science of science, extending the Knowledge Lab study I contributed to: paper novelty labelled by a language model from abstracts with a human-coded calibration set, used to estimate how exposure to an AI breakthrough changes the direction of a scientist’s subsequent work, subject to the relevant data agreements.

7. Timeline and expected contributions

Year 1: coursework and qualifying examinations; Aim 1 theory and the envelope solver; the semi-synthetic testbed. **Year 2:** Aim 2 theory and the clinical transport application; first software release covering Aims 1–2; two papers submitted. **Year 3:** Aim 3 theory, the abstention framework, and the clinical text application. **Year 4:** science-of-science application, integration, and thesis. Each aim yields a self-contained paper, and Aims 2 and 3 do not depend on Aim 1. The proposal will deliver design-sensitivity theory and optimal contrasts for covariate-adaptive full matching; a two-parameter sensitivity analysis that treats hidden bias and population shift jointly; design-based tests for model-generated outcomes and covariates whose validity does not rest on model accuracy; a principled abstention rule; and open-source software in R and Python. The programme sits at the intersection of the research of Prof. Dongming Huang, Prof. Yao Zhang, and Prof. Doudou Zhou, and I would be glad to develop it under their supervision.

References

- [1] Paul R. Rosenbaum. *Observational Studies*. Springer, 2nd edition, 2002.
- [2] **Qi Huang***, Pengyun Wang*, and Xinran Li. Covariate-adaptive studentized sensitivity analysis for Neyman’s weak null in matched designs. Working paper, 2026. (*equal contribution)
- [3] Jason Wu and Peng Ding. Randomization tests for weak null hypotheses in randomized experiments. *Journal of the American Statistical Association*, 116(536):1898–1913, 2021.

- [4] Colin B. Fogarty. Studentized sensitivity analysis for the sample average treatment effect in paired observational studies. *Journal of the American Statistical Association*, 115(531):1518–1530, 2020.
- [5] Kevin Guo and Dominik Rothenhäusler. On the statistical role of inexact matching in observational studies. *Biometrika*, 110(3):631–644, 2023.
- [6] Samuel D. Pimentel and Yaxuan Huang. Covariate-adaptive randomization inference in matched designs. *Journal of the Royal Statistical Society: Series B*, 86(5):1312–1338, 2024.
- [7] Ben B. Hansen. Full matching in an observational study of coaching for the SAT. *Journal of the American Statistical Association*, 99(467):609–618, 2004.
- [8] Paul R. Rosenbaum. Design sensitivity and efficiency in observational studies. *Journal of the American Statistical Association*, 105(490):692–702, 2010.
- [9] Ruth Heller, Paul R. Rosenbaum, and Dylan S. Small. Split samples and design sensitivity in observational studies. *Journal of the American Statistical Association*, 104(487):1090–1101, 2009.
- [10] Elizabeth A. Stuart, Stephen R. Cole, Catherine P. Bradshaw, and Philip J. Leaf. The use of propensity scores to assess the generalizability of results from randomized trials. *Journal of the Royal Statistical Society: Series A*, 174(2):369–386, 2011.
- [11] Issa J. Dahabreh, Sarah E. Robertson, Jon A. Steingrimsso, Elizabeth A. Stuart, and Miguel A. Hernán. Extending inferences from a randomized trial to a new target population. *Statistics in Medicine*, 39(14):1999–2014, 2020.
- [12] Trang Quynh Nguyen, Cyrus Ebnesajjad, Stephen R. Cole, and Elizabeth A. Stuart. Sensitivity analysis for an unobserved moderator in RCT-to-target-population generalization of treatment effects. *Annals of Applied Statistics*, 11(1):225–247, 2017.
- [13] Melody Y. Huang. Sensitivity analysis for the generalization of experimental results. *Journal of the Royal Statistical Society: Series A*, 187(4):900–918, 2024.
- [14] Qingyuan Zhao, Dylan S. Small, and Bhaswar B. Bhattacharya. Sensitivity analysis for inverse probability weighting estimators via the percentile bootstrap. *Journal of the Royal Statistical Society: Series B*, 81(4):735–761, 2019.
- [15] Anastasios N. Angelopoulos, Stephen Bates, Clara Fannjiang, Michael I. Jordan, and Tijana Zrnic. Prediction-powered inference. *Science*, 382(6671):669–674, 2023.
- [16] Anastasios N. Angelopoulos, John C. Duchi, and Tijana Zrnic. PPI++: Efficient prediction-powered inference. arXiv preprint arXiv:2311.01453, 2023.
- [17] Tijana Zrnic and Emmanuel J. Candès. Cross-prediction-powered inference. *Proceedings of the National Academy of Sciences*, 121(15):e2322083121, 2024.
- [18] Naoki Egami, Musashi Hinck, Brandon M. Stewart, and Hanying Wei. Using imperfect surrogates for downstream inference: Design-based supervised learning for social science applications of large language models. In *Advances in Neural Information Processing Systems*, 2023.
- [19] Fangnan Zheng and Yao Zhang. Fit CATE once: Model-assisted randomization tests without sample splitting. arXiv preprint arXiv:2605.09116, 2026.
- [20] Yao Zhang and Qingyuan Zhao. An average-case sensitivity analysis for unmeasured confounding. *Biometrika*, 113(2):asag030, 2026.
- [21] Richard K. Crump, V. Joseph Hotz, Guido W. Imbens, and Oscar A. Mitnik. Dealing with limited overlap in estimation of average treatment effects. *Biometrika*, 96(1):187–199, 2009.
- [22] Doudou Zhou, Molei Liu, Mengyan Li, and Tianxi Cai. Doubly robust augmented model accuracy transfer inference with high dimensional features. *Journal of the American Statistical Association*, 120(549):524–534, 2025.
- [23] Dongming Huang and Lucas Janson. Relaxing the assumptions of knockoffs by conditioning. *The Annals of Statistics*, 48(5):3021–3042, 2020.
- [24] Xiaotong Lin, Weihao Li, Fangqiao Tian, and Dongming Huang. Goodness-of-fit tests for high-dimensional Gaussian graphical models via exchangeable sampling. *Journal of the Royal Statistical Society: Series B*, 87(5):1619–1635, 2025.
- [25] Dongming Huang, Feicheng Wang, Donald B. Rubin, and S. C. Kou. Catalytic priors: Using synthetic data to specify prior distributions in Bayesian analysis. *Statistical Science*, 40(2):259–271, 2025.
- [26] Yao Zhang and Emmanuel J. Candès. Posterior conformal prediction. arXiv preprint arXiv:2409.19712, 2024.
- [27] Zhi Zhang, Lingfeng Lyu, Yue Kang, and Doudou Zhou. Conditional evaluation of language models with cheap auxiliary signals. arXiv preprint arXiv:2608.16210, 2026.
- [28] Longtian Shi, Molei Liu, and Doudou Zhou. Inferential evaluation of surrogate-derived models under covariate shift. arXiv preprint arXiv:2608.15783, 2026.